

# Miller–Rabin Primality Testing: Analysis and Benchmarks

Course Project

December 1, 2025

## Abstract

We present a rigorous analysis of the Miller–Rabin (MR) primality test, its error bounds, and time complexity, alongside benchmarks comparing trial division, Fermat, randomized MR with  $k$  bases, and deterministic MR for 64-bit integers. Experiments span multiple input distributions (random odd, Carmichael, small-factor composites, primes), and we include a toy RSA demo built on the same MR implementation.

## 1 Problem Setting and Model

We seek to decide primality of an integer  $n \geq 2$ . The analysis is stated in the uniform RAM model with unit-cost operations on  $\Theta(\log n)$ -bit words. No implementation details appear in this section; they are deferred to later benchmarking notes.

## 2 Miller–Rabin Algorithm and Proofs

### 2.1 Preprocessing for odd inputs

For every odd  $n > 2$ , factor

$$n - 1 = 2^s \cdot d, \quad d \text{ odd}, \quad s \geq 1.$$

The pair  $(s, d)$  is computed once and reused across all bases.

### 2.2 One Miller–Rabin round

Let  $n > 2$  be odd and  $a$  satisfy  $2 \leq a \leq n - 2$  and  $\gcd(a, n) = 1$ .

1. Compute  $x = a^d \pmod{n}$  by repeated squaring.
2. If  $x \in \{1, n - 1\}$ , accept the round.
3. Otherwise, for  $r = 1, \dots, s - 1$ :
  - (a) Set  $x \leftarrow x^2 \pmod{n}$ .
  - (b) If  $x = n - 1$ , accept; if  $x = 1$ , reject (non-trivial square root of 1).
4. If no acceptance occurs, reject.

### 2.3 Correctness on prime inputs

Assume  $n$  is prime. The multiplicative group  $G = (\mathbb{Z}/n\mathbb{Z})^\times$  has order  $|G| = n - 1 = 2^s d$ .

**Lemma 1** (Square roots of unity). *If  $x^2 \equiv 1 \pmod{n}$ , then  $x \equiv \pm 1 \pmod{n}$ .*

*Proof.* We have  $n \mid (x - 1)(x + 1)$ . Since  $n$  is prime, one factor vanishes modulo  $n$ .  $\square$

**Lemma 2** (Allowed trajectory). *Let  $x = a^d \in G$ . Then  $x^{2^s} = a^{n-1} \equiv 1$ . In the sequence  $x, x^2, x^4, \dots, x^{2^s}$ , any transition from a value other than  $-1$  to  $1$  would create a non-trivial square root of  $1$ , contradicting Lemma 1. Therefore the sequence must either start at  $1$  or hit  $-1$  exactly once before reaching  $1$ .*

*Proof.* The equality  $x^{2^s} = 1$  follows from Lagrange's theorem. Suppose some iterate  $x^{2^j} = 1$  with  $0 < j < s$  and  $x^{2^{j-1}} \not\equiv -1$ . Then  $x^{2^{j-1}}$  is a non-trivial square root of  $1$ , impossible by Lemma 1. Hence acceptance in a prime modulus always occurs along the Miller–Rabin path.  $\square$

**Theorem 1** (No false negatives). *For prime  $n$  and any admissible base  $a$ , one Miller–Rabin round accepts.*

*Proof.* Lemma 2 shows the round's control flow accepts exactly in the situations mandated by the group structure; no rejection is possible.  $\square$

## 2.4 Error bound on composite inputs

Let  $n$  be odd and composite; write  $n - 1 = 2^s d$  with  $d$  odd. Define a base  $a$  coprime to  $n$  as a *strong liar* if the round accepts and as a *strong witness* otherwise. We prove that at most a quarter of all admissible bases are liars, yielding the classical  $(1/4)^k$  error bound for  $k$  independent rounds.

**CRT decomposition.** Factor  $n = \prod_{i=1}^r p_i^{e_i}$ . By the Chinese Remainder Theorem,

$$(\mathbb{Z}/n\mathbb{Z})^\times \cong \prod_{i=1}^r G_i, \quad G_i = (\mathbb{Z}/p_i^{e_i}\mathbb{Z})^\times.$$

Each  $G_i$  is cyclic of order  $p_i^{e_i-1}(p_i - 1) = 2^{t_i} m_i$  with  $m_i$  odd and  $t_i \geq 1$ . Choose generators  $g_i$  so every  $a$  corresponds to exponents  $u_i$  with  $0 \leq u_i < 2^{t_i} m_i$ .

**Component-wise acceptance.** Write  $u_i = 2^{\lambda_i} v_i$  with  $v_i$  odd and  $0 \leq \lambda_i \leq t_i$ . The MR acceptance condition in component  $i$  is:

$$\text{either } g_i^{u_i d} \equiv 1 \pmod{p_i^{e_i}} \quad \text{or} \quad g_i^{u_i 2^j d} \equiv -1 \text{ for some } j < t_i.$$

Because  $d$  is odd,  $g_i^{u_i d} \equiv 1$  iff  $m_i \mid u_i$ . Moreover,  $g_i^{u_i 2^j d} \equiv -1$  requires  $m_i \mid u_i$  and  $u_i/m_i \equiv 2^{t_i-1} \pmod{2^{t_i}}$ , the unique exponent producing order 2 after the squaring ladder.

**Counting liars per component.** Acceptable  $u_i$  are exactly those with  $m_i \mid u_i$  and  $u_i \pmod{2^{t_i}} \in \{0, 2^{t_i-1}\}$ . The count is

$$|L_i| = 2 \cdot \frac{2^{t_i} m_i}{2^{t_i}} = \frac{|G_i|}{2^{t_i-1}},$$

so the liar fraction in component  $i$  is  $\rho_i = 2^{-(t_i-1)} \leq \frac{1}{2}$ , with  $\rho_i \leq \frac{1}{4}$  whenever  $t_i \geq 3$ .

**Global product bound.** Independence of the CRT components gives  $|L| = \prod_i |L_i|$  and

$$\frac{|L|}{|G|} = \prod_{i=1}^r \rho_i = 2^{-\sum_i(t_i-1)}.$$

For any odd composite  $n$ ,  $\sum_i(t_i - 1) \geq 2$ :

- With at least two distinct prime factors ( $r \geq 2$ ), two components each contribute a factor at most  $\frac{1}{2}$ , yielding  $|L|/|G| \leq \frac{1}{4}$ .
- For a prime power  $p^e$  with  $e \geq 2$ , the 2-adic valuation of  $|G_1|$  satisfies  $t_1 \geq 2$ , so  $\rho_1 \leq \frac{1}{4}$ .

Thus  $|L| \leq \frac{1}{4}|G|$ . Random bases are chosen uniformly from admissible residues, so one round accepts a composite with probability at most  $\frac{1}{4}$ , and  $k$  independent rounds accept with probability at most  $(1/4)^k$ .

## 2.5 Deterministic base set for 64-bit inputs

The fixed bases  $\{2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31, 37\}$  suffice for all odd  $n < 2^{64}$ : passing MR for all of them implies primality. Therefore `is_prime_mr_det64` is exact on 64-bit integers.

## 2.6 Time Complexity

Let  $M(n)$  denote the cost of one modular multiplication on  $\Theta(\log n)$ -bit words; in the RAM model  $M(n) = O(1)$ .

- Fast exponentiation of  $a^d \bmod n$  uses  $\lfloor \log_2 d \rfloor + \text{popcount}(d)$  multiplications, i.e.,  $O(\log n)$ .
- One MR round consists of that exponentiation and up to  $s - 1 \leq \log_2(n - 1)$  additional squarings, so  $O(\log n)$ .
- $k$  rounds take  $O(k \log n)$  time and  $O(1)$  space.
- Fermat- $k$  matches these time bounds but lacks the  $4^{-k}$  error decay; trial division is  $O(\sqrt{n})$ .

## 2.7 Summary of Bounds

Algorithm	Time	Error	Space
Trial Division	$O(\sqrt{n})$	0	$O(1)$
Fermat- $k$	$O(k \log n)$	Can be 1 on Carmichael	$O(1)$
MR- $k$	$O(k \log n)$	$\leq (1/4)^k$	$O(1)$
MR-det-64	$O(\log n)$	0 (for $n < 2^{64}$ )	$O(1)$

Table 1: Asymptotic behavior of implemented algorithms.

## 3 Algorithms Implemented (C++)

- **Trial Division (TD)** (`is_prime_td`): odd divisors up to  $\lfloor \sqrt{n} \rfloor$ ; disabled for bits > 48.
- **Fermat- $k$**  (`is_probable_prime_fermat`):  $k$  random bases; rejects on any violation of  $a^{n-1} \equiv 1$ .
- **MR- $k$**  (`is_probable_prime_mr`):  $k$  random bases with the standard strong probable prime check; per-round error  $\leq 1/4$ , total  $\leq 4^{-k}$ .
- **MR-det-64** (`is_prime_mr_det64`): fixed 12 bases; exact for  $n < 2^{64}$ .

## 4 Benchmark Workflow

All benchmarks live in `src/main.cpp`; `scripts/run_benchmarks.py` drives the binary `mr_bench`.

- **CLI:** `./mr_bench <algo_id> <dist_id> <bits> <sample_count> <rounds> <seed_base> <reps>`.
- **Distributions (dist\_id):** random odd, Carmichael, composite with small factor, primes (generated via MR-det-64).
- **Metrics:** each repetition outputs `<time_ns_total> <error_count>` across the sample of size  $S$ , with truth labels provided by MR-det-64.
- **Rounds:** Fermat/MR use  $k$  supplied on the CLI (e.g.,  $k = 10$ ); TD ignores it.
- **Python aggregation:** `results/raw_results.csv` records all rows; plots are generated by `plot_results.py`.

## 5 Results

### 5.1 Plots

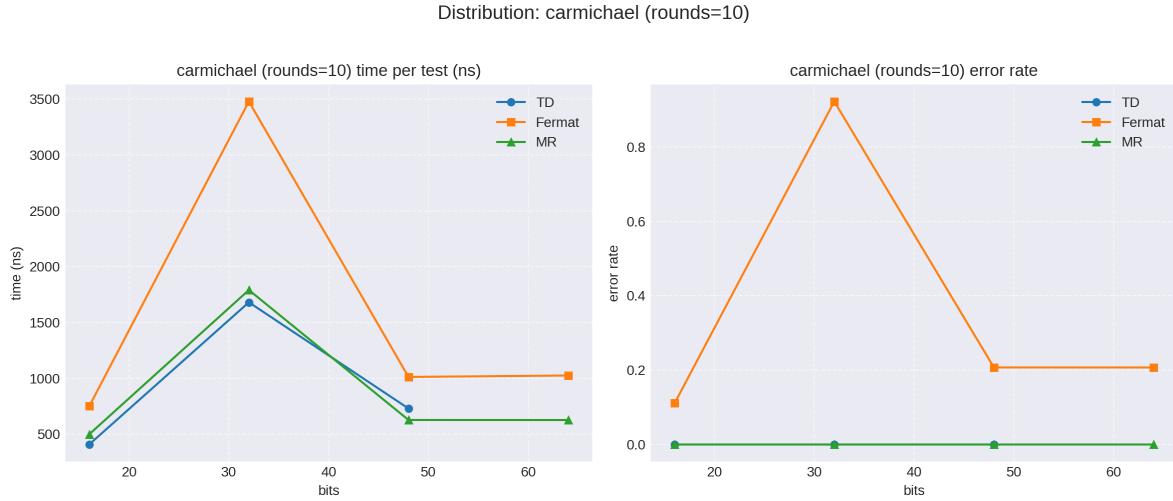


Figure 1: Carmichael distribution,  $k = 10$ : time per test (ns) and error rate.

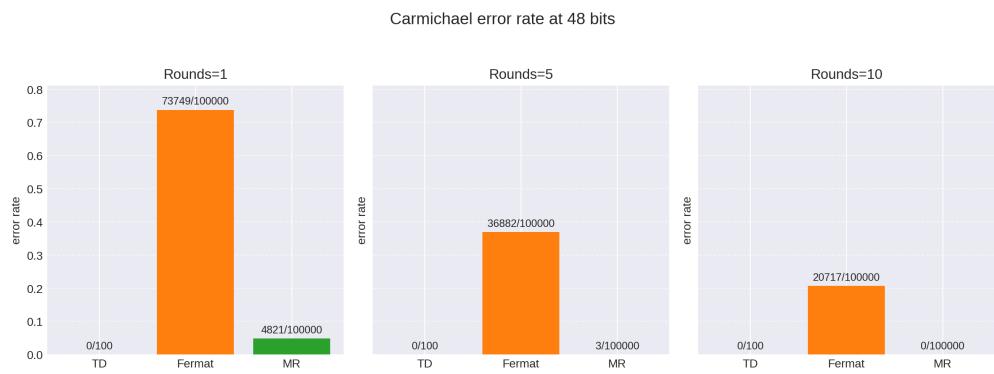


Figure 2: Carmichael, 48 bits: error rate vs rounds for Fermat and MR.

Distribution: rand\_odd (rounds=10)

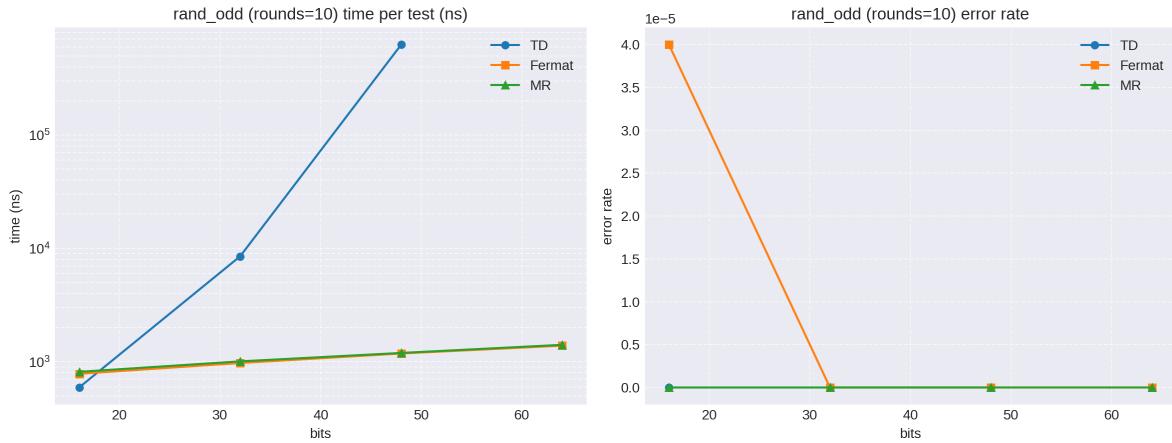


Figure 3: Random odd inputs,  $k = 10$ : time per test (ns) and error rate.

Distribution: primes (rounds=10)

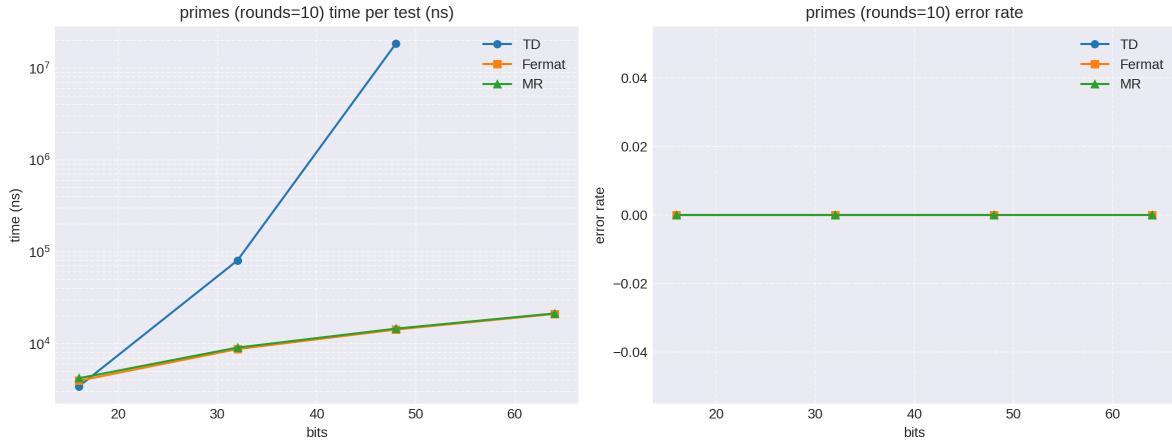


Figure 4: Prime inputs (ground truth),  $k = 10$ : time per test (ns) and false-negative rate.

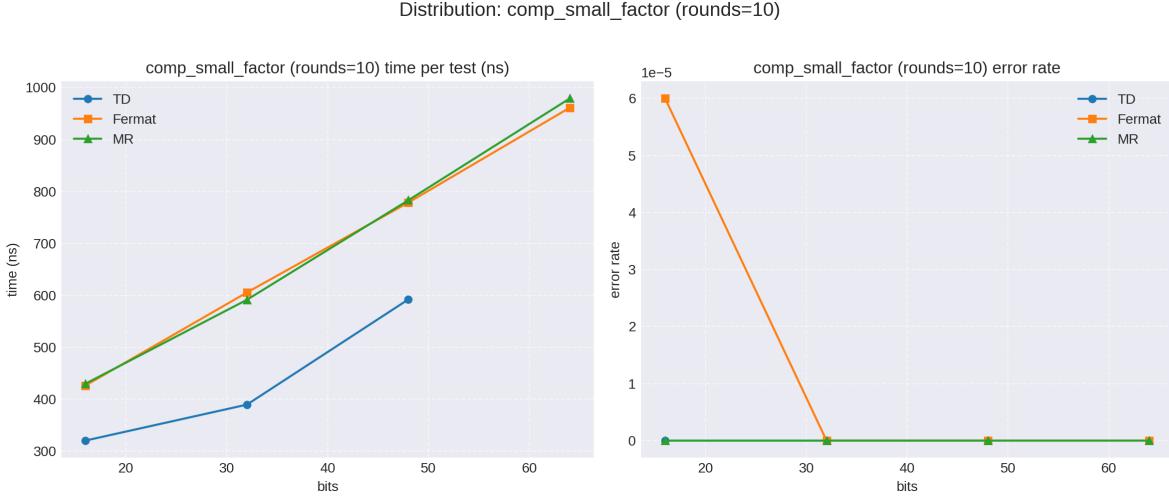


Figure 5: Composites with small factor,  $k = 10$ : time per test (ns) and error rate.

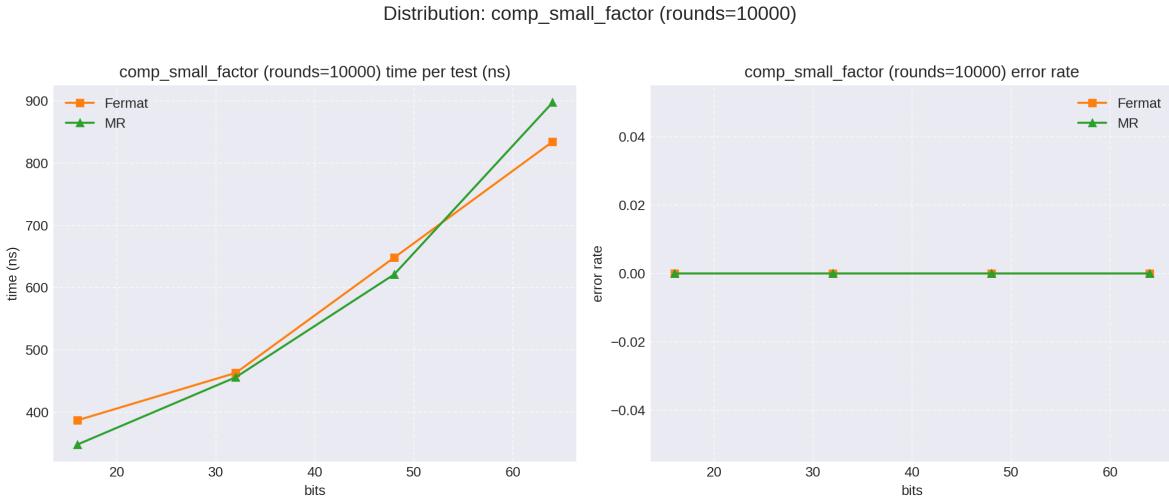


Figure 6: Composites with small factor,  $k = 10000$ : time per test (ns) and error rate.

## 5.2 Observations

- **Carmichael sensitivity:** Figure 1 shows MR attaining near-zero error on Carmichael inputs at  $k = 10$ ; Figure 2 highlights the exponential decay, with error essentially gone by  $k = 5$ .
- **Random odd mix:** In Figure 3, Fermat and MR have similar runtimes; MR keeps error negligible, and Fermat's empirical error is extremely rare (about  $4 \times 10^{-5}$ ).
- **Prime inputs:** Figure 4 shows no false negatives for MR (as guaranteed); Fermat also very rarely rejects primes.
- **Small-factor composites:** Figures 5 and 6 demonstrate that MR's error vanishes quickly with  $k$ ; time scales linearly with  $k$ , matching the  $O(k \log n)$  or  $O(k \log^2 n)$  bounds depending on the cost model.

## 6 Real-world Usage

- **RSA key generation:** OpenSSL applies multiple MR rounds (and sieving) to candidate primes; documented error probabilities are bounded by powers of two. The toy RSA demo in this repository reuses the same MR code for small key generation.
- **Big-integer libraries:** GMP, Java `BigInteger.isProbablePrime`, and similar libraries rely on MR (often combined with Lucas/Baillie-PSW) to deliver negligible error on large inputs.
- **Computer algebra systems:** CAS tools (Maple, Mathematica, PARI/GP, SageMath) embed MR-style strong probable prime tests as core building blocks.
- **Toy RSA implementation:** A pedagogical RSA demo (local CLI and simple client/server) is included; it generates small primes using the MR routines documented here and walks through textbook RSA key generation, encryption, and decryption for illustration.

## 7 Conclusions

Miller–Rabin delivers logarithmic-time probabilistic primality testing with exponentially small error in the number of rounds. Deterministic base sets make it exact for 64-bit integers, and randomized MR dominates Fermat, especially on adversarial inputs like Carmichael numbers. Benchmarks confirm the theoretical bounds: time grows linearly with  $k \log n$ , and empirical error aligns with the  $(1/4)^k$  guarantee.